

```
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.layers import Embedding, LSTM, Dense
from tensorflow.keras.models import Sequential
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers import Adam
import pickle
import numpy as np
import os
```

```
from google.colab import files
uploaded = files.upload()
```

No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving Pride and Prejudice.txt to Pride and Prejudice (3).txt

```
file = open("Pride and Prejudice.txt", "r", encoding = "utf8")

# store file in list
lines = []
for i in file:
    lines.append(i)

# Convert list to string
data = ""
for i in lines:
    data = ' '.join(lines)

#replace unnecessary stuff with space
data = data.replace('\n', ' ').replace('\r', ' ').replace('\uffff', ' ').replace('"','').replace("'",') #new line, carriage return, unicode ch

#remove unnecessary spaces
data = data.split()
data = ' '.join(data)
data[:500]
```

☞ 'The Project Gutenberg eBook of Pride and prejudice, by Jane Austen This eBook is for the use of anyone anywhere in the United States and most other parts of the world at no cost and with almost no restrictions whatsoever. You may copy it, give it away or re-use it under the terms of the Project Gutenberg License included with this eBook or online at www.gutenberg.org. If you are not located in the United

```
len(data)
```

```
733851
```

```
tokenizer = Tokenizer()
tokenizer.fit_on_texts([data])
```

```
# saving the tokenizer for predict function
pickle.dump(tokenizer, open('token.pkl', 'wb'))
```

```
sequence_data = tokenizer.texts_to_sequences([data])[0]
sequence_data[:15]
```

```
[1, 182, 164, 1001, 3, 299, 4, 946, 30, 72, 710, 41, 1001, 23, 21]
```

```
len(sequence_data)
```

```
131237
```

```
vocab_size = len(tokenizer.word_index) + 1
print(vocab_size)
```

```
7250
```

```

sequences = []

for i in range(3, len(sequence_data)):
    words = sequence_data[i-3:i+1]
    sequences.append(words)

print("The Length of sequences are: ", len(sequences))
sequences = np.array(sequences)
sequences[:10]

```

```

The Length of sequences are: 131234
array([[ 1, 182, 164, 1001],
       [182, 164, 1001, 3],
       [164, 1001, 3, 299],
       [1001, 3, 299, 4],
       [ 3, 299, 4, 946],
       [299, 4, 946, 30],
       [ 4, 946, 30, 72],
       [946, 30, 72, 710],
       [ 30, 72, 710, 41],
       [ 72, 710, 41, 1001]])

```

```

X = []
y = []

for i in sequences:
    X.append(i[0:3])
    y.append(i[3])

```

```

X = np.array(X)
y = np.array(y)

```

```

print("Data: ", X[:10])
print("Response: ", y[:10])

```

```

Data: [[ 1 182 164]
       [182 164 1001]
       [164 1001 3]
       [1001 3 299]
       [ 3 299 4]
       [299 4 946]
       [ 4 946 30]
       [946 30 72]
       [ 30 72 710]
       [ 72 710 41]]
Response: [1001 3 299 4 946 30 72 710 41 1001]

```

```

y = to_categorical(y, num_classes=vocab_size)
y[:5]

```

```

array([[0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.]], dtype=float32)

```

```

model = Sequential()
model.add(Embedding(vocab_size, 10, input_length=3))
model.add(LSTM(1000, return_sequences=True))
model.add(LSTM(1000))
model.add(Dense(1000, activation="relu"))
model.add(Dense(vocab_size, activation="softmax"))

```

```
model.summary()
```

```
Model: "sequential"
```


Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 3, 10)	72500

lstm (LSTM)	(None, 3, 1000)	4044000
lstm_1 (LSTM)	(None, 1000)	8004000
dense (Dense)	(None, 1000)	1001000
dense_1 (Dense)	(None, 7250)	7257250

```
=====
Total params: 20,378,750
Trainable params: 20,378,750
Non-trainable params: 0
=====
```

```
from tensorflow import keras
from keras.utils.vis_utils import plot_model

keras.utils.plot_model(model, to_file='plot.png', show_layer_names=True)
```



```
from tensorflow.keras.callbacks import ModelCheckpoint

checkpoint = ModelCheckpoint("next_words.h5", monitor='loss', verbose=1, save_best_only=True)
model.compile(loss="categorical_crossentropy", optimizer=Adam(learning_rate=0.001))
model.fit(X, y, epochs=70, batch_size=64, callbacks=[checkpoint])
```

```

Epoch 37: loss improved from 0.63283 to 0.62363, saving model to next_words.h5
2051/2051 [=====] - 33s 16ms/step - loss: 0.6236
Epoch 38/70
2049/2051 [=====>.] - ETA: 0s - loss: 0.6088
Epoch 38: loss improved from 0.62363 to 0.60894, saving model to next_words.h5
2051/2051 [=====] - 33s 16ms/step - loss: 0.6089
Epoch 39/70
2051/2051 [=====] - ETA: 0s - loss: 0.6003
Epoch 39: loss improved from 0.60894 to 0.60032, saving model to next_words.h5
2051/2051 [=====] - 33s 16ms/step - loss: 0.6003
Epoch 40/70
2049/2051 [=====>.] - ETA: 0s - loss: 0.5863
Epoch 40: loss improved from 0.60032 to 0.58638, saving model to next_words.h5
2051/2051 [=====] - 36s 18ms/step - loss: 0.5864
Epoch 41/70
2050/2051 [=====>.] - ETA: 0s - loss: 0.5787
Epoch 41: loss improved from 0.58638 to 0.57879, saving model to next_words.h5
2051/2051 [=====] - 32s 15ms/step - loss: 0.5788
Epoch 42/70
2051/2051 [=====] - ETA: 0s - loss: 0.5727
Epoch 42: loss improved from 0.57879 to 0.57266, saving model to next_words.h5
2051/2051 [=====] - 31s 15ms/step - loss: 0.5727
Epoch 43/70
2048/2051 [=====>.] - ETA: 0s - loss: 0.5607
Epoch 43: loss improved from 0.57266 to 0.56090, saving model to next_words.h5
2051/2051 [=====] - 34s 16ms/step - loss: 0.5609
Epoch 44/70
2048/2051 [=====>.] - ETA: 0s - loss: 0.5555
Epoch 44: loss improved from 0.56090 to 0.55557, saving model to next_words.h5
2051/2051 [=====] - 35s 17ms/step - loss: 0.5556
Epoch 45/70
2049/2051 [=====>.] - ETA: 0s - loss: 0.5483
Epoch 45: loss improved from 0.55557 to 0.54833, saving model to next_words.h5
2051/2051 [=====] - 34s 16ms/step - loss: 0.5483
Epoch 46/70
2050/2051 [=====>.] - ETA: 0s - loss: 0.5410
Epoch 46: loss improved from 0.54833 to 0.54105, saving model to next_words.h5

```

```

import joblib
filename = 'next_words.h5'
joblib.dump(model, filename)

```

```
['next_words.h5']
```

```

from tensorflow.keras.models import load_model
import numpy as np
import pickle

# Load the model and tokenizer
model = load_model('next_words.h5')
tokenizer = pickle.load(open('token.pkl', 'rb'))

def Predict_Next_Words(model, tokenizer, text):

    sequence = tokenizer.texts_to_sequences([text])
    sequence = np.array(sequence)
    preds = np.argmax(model.predict(sequence))
    predicted_word = ""

    for key, value in tokenizer.word_index.items():
        if value == preds:
            predicted_word = key
            break

    print(predicted_word)
    return predicted_word

while(True):
    text = input("Enter your line: ")

    if text == "0":
        print("Execution completed....")
        break

    else:
        try:
            text = text.split(" ")

```

```
text = text[-3:]
print(text)
```

```
Predict_Next_Words(model, tokenizer, text)
```

```
except Exception as e:
    print("Error occurred: ",e)
    continue
```

```
Enter your line: give it away
['give', 'it', 'away']
1/1 [=====] - 1s 752ms/step
or
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-17-a124736db0da> in <cell line: 1>()
      1 while(True):
----> 2     text = input("Enter your line: ")
      3
      4     if text == "0":
      5         print("Execution completed....")
```

1 frames

```
/usr/local/lib/python3.10/dist-packages/ipykernel/kernelbase.py in
_input_request(self, prompt, ident, parent, password)
    893         except KeyboardInterrupt:
    894             # re-raise KeyboardInterrupt, to truncate traceback
--> 895             raise KeyboardInterrupt("Interrupted by user") from None
    896         except Exception as e:
    897             self.log.warning("Invalid Message:", exc_info=True)
```

KeyboardInterrupt: Interrupted by user

SEARCH STACK OVERFLOW

```
%%writefile app.py
import streamlit as st
import pandas as pd
import numpy as np
import joblib
```

```
st.title('Next Word Predictor')
```

```
S0 = st.number_input("Enter the word:")
loaded_model = joblib.load('next_words.h5')
inputs = (S0)
if st.button("Predict"):
    result = loaded_model.predict([inputs])
    st.write(result)
```

Overwriting app.py

```
!ngrok authtoken 2Kgk0JPvh53BRHvj2SCxcZI3YPN_2eDZG97kbtqGH7xy3yD88
```

Authtoken saved to configuration file: /root/.ngrok2/ngrok.yml

```
!wget https://bin.equinox.io/c/4VmDzA7iaHb/ngrok-stable-linux-amd64.zip
```

```
--2023-07-06 19:04:32-- https://bin.equinox.io/c/4VmDzA7iaHb/ngrok-stable-linux-amd64.zip
Resolving bin.equinox.io (bin.equinox.io)... 52.202.168.65, 54.237.133.81, 18.205.222.128, ...
Connecting to bin.equinox.io (bin.equinox.io)|52.202.168.65|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 13921656 (13M) [application/octet-stream]
Saving to: 'ngrok-stable-linux-amd64.zip.3'
```

```
ngrok-stable-linux- 100%[=====] 13.28M 13.9MB/s in 1.0s
```

```
2023-07-06 19:04:33 (13.9 MB/s) - 'ngrok-stable-linux-amd64.zip.3' saved [13921656/13921656]
```

```
!unzip ngrok-stable-linux-amd64.zip
```

```
Archive: ngrok-stable-linux-amd64.zip
replace ngrok? [y]es, [n]o, [A]ll, [N]one, [r]ename: A
inflating: ngrok
```

```
get_ipython().system_raw('./ngrok http 8501 &')
```

```
! curl -s http://localhost:4040/api/tunnels | python3 -c \
"import sys, json; print(json.load(sys.stdin)['tunnels'][0]['public_url'])"
```

<https://d2a7-34-90-196-79.ngrok-free.app>

```
!streamlit run /content/app.py
```

Collecting usage statistics. To deactivate, set browser.gatherUsageStats to False.

You can now view your Streamlit app in your browser.

Network URL: <http://172.28.0.12:8501>
External URL: <http://34.90.196.79:8501>

Stopping...
Stopping...